

The flower dataset was downloaded from the link

Out of the images, 42 daisy images and 28 dandelion images were chosen for training

7 daisy and 8 dandelion images were chosen for testing, due to time constraint and limited computational ability

Labelme was used to annotate the dataset and the annotation for each file was exported as a json file





```
In [1]: #importing Libraries
from detectron2.engine import DefaultTrainer
from detectron2.config import get_cfg
import detectron2
from detectron2.utils.logger import setup_logger
setup_logger()

import numpy as np
import os, json, cv2, random

from detectron2 import model_zoo
from detectron2.engine import DefaultPredictor
from detectron2.config import get_cfg
from detectron2.utils.visualizer import Visualizer
from detectron2.data import MetadataCatalog, DatasetCatalog
from detectron2.structures import BoxMode
import os
import numpy as np
import json
import random
import matplotlib.pyplot as plt
%matplotlib inline

from detectron2.structures import BoxMode
from detectron2.data import DatasetCatalog, MetadataCatalog
import torch
#assert torch.__version__.startswith("1.8")
import torchvision
import cv2
from detectron2.engine import DefaultTrainer
```

```
C:\Users\shash\anaconda3\envs\detect\lib\site-packages\tqdm\auto.py:22: TqdmWarning:
IPProgress not found. Please update jupyter and ipywidgets. See https://ipywidgets.readthedocs.io/en/stable/user_install.html
  from .autonotebook import tqdm as notebook_tqdm
```

```
In [2]: ##define classes
classes = ['daisy', 'dandelion']
```

```
In [3]: ##function for uploading a dataset to detectron2##
def get_data_dicts(directory, classes):
    dataset_dicts = []
    for filename in [file for file in os.listdir(directory) if file.endswith('.json')]:
        json_file = os.path.join(directory, filename)
```

```

with open(json_file) as f:
    img_anns = json.load(f)

record = {}

filename = os.path.join(directory, img_anns["imagePath"])

record["file_name"] = filename
record["height"] = 512
record["width"] = 512

annos = img_anns["shapes"]
objs = []
for anno in annos:
    px = [a[0] for a in anno['points']] # x coord
    py = [a[1] for a in anno['points']] # y-coord
    poly = [(x, y) for x, y in zip(px, py)] # poly for segmentation
    poly = [p for x in poly for p in x]

    obj = {
        "bbox": [np.min(px), np.min(py), np.max(px), np.max(py)],
        "bbox_mode": BoxMode.XYXY_ABS,
        "segmentation": [poly],
        "category_id": classes.index(anno['label']),
        "iscrowd": 0
    }
    objs.append(obj)
record["annotations"] = objs
dataset_dicts.append(record)
return dataset_dicts

```

```

In [4]: #uploading data
for d in ["train", "test"]:
    DatasetCatalog.register(
        "outlierz" + d,
        lambda d=d: get_data_dicts(d, classes)
    )
    MetadataCatalog.get("outlierz" + d).set(thing_classes=classes)

```

```

In [6]: #saving metadata
microcontroller_metadata = MetadataCatalog.get("outlierztrain")

```

```

In [7]: cfg = get_cfg()

cfg.merge_from_file(model_zoo.get_config_file("COCO-InstanceSegmentation/mask_rcnn_R_50_FPN_3x.yaml"))
cfg.DATASETS.TRAIN = ("outlierztrain",)
cfg.DATASETS.TEST = ()
cfg.DATALOADER.NUM_WORKERS = 0
cfg.MODEL.WEIGHTS = model_zoo.get_checkpoint_url("COCO-InstanceSegmentation/mask_rcnn_R_50_FPN_3x.yaml")
# Let training initialize from model zoo
cfg.SOLVER.IMS_PER_BATCH = 2
cfg.SOLVER.BASE_LR = 0.00025
# pick a good LR
cfg.SOLVER.MAX_ITER = 300
# 300 iterations seems good enough for this toy dataset; you will need to train Longer
cfg.SOLVER.STEPS = []
# do not decay Learning rate
cfg.MODEL.ROI_HEADS.BATCH_SIZE_PER_IMAGE = 128

```

```
# faster, and good enough for this toy dataset (default: 512)
cfg.MODEL.ROI_HEADS.NUM_CLASSES = 2
# only has one class (balloon). (see https://detectron2.readthedocs.io/tutorials/datasets.html)
# NOTE: this config means the number of classes, but a few popular unofficial tutorials
# use 2

os.makedirs(cfg.OUTPUT_DIR, exist_ok=True)
trainer = DefaultTrainer(cfg)
trainer.resume_or_load(resume=False)
trainer.train()
```

```
[10/09 12:23:01 d2.engine.defaults]: Model:
GeneralizedRCNN(
    (backbone): FPN(
        (fpn_lateral2): Conv2d(256, 256, kernel_size=(1, 1), stride=(1, 1))
        (fpn_output2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
        (fpn_lateral3): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1))
        (fpn_output3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
        (fpn_lateral4): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1))
        (fpn_output4): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
        (fpn_lateral5): Conv2d(2048, 256, kernel_size=(1, 1), stride=(1, 1))
        (fpn_output5): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
    )
    (top_block): LastLevelMaxPool()
    (bottom_up): ResNet(
        (stem): BasicStem(
            (conv1): Conv2d(
                3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False
                (norm): FrozenBatchNorm2d(num_features=64, eps=1e-05)
            )
        )
        (res2): Sequential(
            (0): BottleneckBlock(
                (shortcut): Conv2d(
                    64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
                    (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
                )
                (conv1): Conv2d(
                    64, 64, kernel_size=(1, 1), stride=(1, 1), bias=False
                    (norm): FrozenBatchNorm2d(num_features=64, eps=1e-05)
                )
                (conv2): Conv2d(
                    64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
                    (norm): FrozenBatchNorm2d(num_features=64, eps=1e-05)
                )
                (conv3): Conv2d(
                    64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
                    (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
                )
            )
            (1): BottleneckBlock(
                (conv1): Conv2d(
                    256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False
                    (norm): FrozenBatchNorm2d(num_features=64, eps=1e-05)
                )
                (conv2): Conv2d(
                    64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
                    (norm): FrozenBatchNorm2d(num_features=64, eps=1e-05)
                )
                (conv3): Conv2d(
                    64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
                    (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
                )
            )
            (2): BottleneckBlock(
                (conv1): Conv2d(
                    256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False
                    (norm): FrozenBatchNorm2d(num_features=64, eps=1e-05)
                )
            )
        )
    )
)
```

```
)  
(conv2): Conv2d(  
    64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False  
    (norm): FrozenBatchNorm2d(num_features=64, eps=1e-05)  
)  
(conv3): Conv2d(  
    64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False  
    (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)  
)  
)  
(res3): Sequential(  
    (0): BottleneckBlock(  
        (shortcut): Conv2d(  
            256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False  
            (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)  
)  
        (conv1): Conv2d(  
            256, 128, kernel_size=(1, 1), stride=(2, 2), bias=False  
            (norm): FrozenBatchNorm2d(num_features=128, eps=1e-05)  
)  
        (conv2): Conv2d(  
            128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False  
            (norm): FrozenBatchNorm2d(num_features=128, eps=1e-05)  
)  
        (conv3): Conv2d(  
            128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False  
            (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)  
)  
)  
(1): BottleneckBlock(  
    (conv1): Conv2d(  
        512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False  
        (norm): FrozenBatchNorm2d(num_features=128, eps=1e-05)  
)  
    (conv2): Conv2d(  
        128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False  
        (norm): FrozenBatchNorm2d(num_features=128, eps=1e-05)  
)  
    (conv3): Conv2d(  
        128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False  
        (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)  
)  
)  
(2): BottleneckBlock(  
    (conv1): Conv2d(  
        512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False  
        (norm): FrozenBatchNorm2d(num_features=128, eps=1e-05)  
)  
    (conv2): Conv2d(  
        128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False  
        (norm): FrozenBatchNorm2d(num_features=128, eps=1e-05)  
)  
    (conv3): Conv2d(  
        128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False  
        (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)  
)  
)  
(3): BottleneckBlock(  
    (conv1): Conv2d(  
        512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False  
        (norm): FrozenBatchNorm2d(num_features=128, eps=1e-05)  
)  
    (conv2): Conv2d(  
        128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False  
        (norm): FrozenBatchNorm2d(num_features=128, eps=1e-05)  
)  
    (conv3): Conv2d(  
        128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False  
        (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)  
)  
)
```

```
512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False
  (norm): FrozenBatchNorm2d(num_features=128, eps=1e-05)
)
(conv2): Conv2d(
  128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
  (norm): FrozenBatchNorm2d(num_features=128, eps=1e-05)
)
(conv3): Conv2d(
  128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False
  (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)
)
)
)
(res4): Sequential(
  (0): BottleneckBlock(
    (shortcut): Conv2d(
      512, 1024, kernel_size=(1, 1), stride=(2, 2), bias=False
      (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
    )
    (conv1): Conv2d(
      512, 256, kernel_size=(1, 1), stride=(2, 2), bias=False
      (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv2): Conv2d(
      256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
      (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv3): Conv2d(
      256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False
      (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
    )
  )
  (1): BottleneckBlock(
    (conv1): Conv2d(
      1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
      (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv2): Conv2d(
      256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
      (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv3): Conv2d(
      256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False
      (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
    )
  )
  (2): BottleneckBlock(
    (conv1): Conv2d(
      1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
      (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv2): Conv2d(
      256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
      (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv3): Conv2d(
      256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False
      (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
    )
  )
)
```

```
(3): BottleneckBlock(
    (conv1): Conv2d(
        1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv2): Conv2d(
        256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv3): Conv2d(
        256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
    )
)
(4): BottleneckBlock(
    (conv1): Conv2d(
        1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv2): Conv2d(
        256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv3): Conv2d(
        256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
    )
)
(5): BottleneckBlock(
    (conv1): Conv2d(
        1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv2): Conv2d(
        256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv3): Conv2d(
        256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
    )
)
)
(res5): Sequential(
    (0): BottleneckBlock(
        (shortcut): Conv2d(
            1024, 2048, kernel_size=(1, 1), stride=(2, 2), bias=False
            (norm): FrozenBatchNorm2d(num_features=2048, eps=1e-05)
        )
        (conv1): Conv2d(
            1024, 512, kernel_size=(1, 1), stride=(2, 2), bias=False
            (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)
        )
        (conv2): Conv2d(
            512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
            (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)
        )
        (conv3): Conv2d(
            512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False
            (norm): FrozenBatchNorm2d(num_features=2048, eps=1e-05)
        )
)
```

```

        )
    )
(1): BottleneckBlock(
    (conv1): Conv2d(
        2048, 512, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)
    )
    (conv2): Conv2d(
        512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)
    )
    (conv3): Conv2d(
        512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=2048, eps=1e-05)
    )
)
(2): BottleneckBlock(
    (conv1): Conv2d(
        2048, 512, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)
    )
    (conv2): Conv2d(
        512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)
    )
    (conv3): Conv2d(
        512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=2048, eps=1e-05)
    )
)
)
)
)
)
(proposal_generator): RPN(
    (rpn_head): StandardRPNHead(
        (conv): Conv2d(
            256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1)
            (activation): ReLU()
        )
        (objectness_logits): Conv2d(256, 3, kernel_size=(1, 1), stride=(1, 1))
        (anchor_deltas): Conv2d(256, 12, kernel_size=(1, 1), stride=(1, 1))
    )
    (anchor_generator): DefaultAnchorGenerator(
        (cell_anchors): BufferList()
    )
)
(roi_heads): StandardROIHeads(
    (box_pooler): ROIPooler(
        (level_poolers): ModuleList(
            (0): ROIAlign(output_size=(7, 7), spatial_scale=0.25, sampling_ratio=0, aligned=True)
            (1): ROIAlign(output_size=(7, 7), spatial_scale=0.125, sampling_ratio=0, aligned=True)
            (2): ROIAlign(output_size=(7, 7), spatial_scale=0.0625, sampling_ratio=0, aligned=True)
            (3): ROIAlign(output_size=(7, 7), spatial_scale=0.03125, sampling_ratio=0, aligned=True)
        )
    )
    (box_head): FastRCNNConvFCHead(

```

```

        (flatten): Flatten(start_dim=1, end_dim=-1)
        (fc1): Linear(in_features=12544, out_features=1024, bias=True)
        (fc_relu1): ReLU()
        (fc2): Linear(in_features=1024, out_features=1024, bias=True)
        (fc_relu2): ReLU()
    )
    (box_predictor): FastRCNNOutputLayers(
        (cls_score): Linear(in_features=1024, out_features=3, bias=True)
        (bbox_pred): Linear(in_features=1024, out_features=8, bias=True)
    )
    (mask_pooler): ROIAligner(
        (level_poolers): ModuleList(
            (0): ROIAlign(output_size=(14, 14), spatial_scale=0.25, sampling_ratio=0, ali
gned=True)
            (1): ROIAlign(output_size=(14, 14), spatial_scale=0.125, sampling_ratio=0, ali
gned=True)
            (2): ROIAlign(output_size=(14, 14), spatial_scale=0.0625, sampling_ratio=0, ali
gned=True)
            (3): ROIAlign(output_size=(14, 14), spatial_scale=0.03125, sampling_ratio=0, ali
gned=True)
        )
    )
    (mask_head): MaskRCNNConvUpsampleHead(
        (mask_fcn1): Conv2d(
            256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1)
            (activation): ReLU()
        )
        (mask_fcn2): Conv2d(
            256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1)
            (activation): ReLU()
        )
        (mask_fcn3): Conv2d(
            256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1)
            (activation): ReLU()
        )
        (mask_fcn4): Conv2d(
            256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1)
            (activation): ReLU()
        )
        (deconv): ConvTranspose2d(256, 256, kernel_size=(2, 2), stride=(2, 2))
        (deconv_relu): ReLU()
        (predictor): Conv2d(256, 2, kernel_size=(1, 1), stride=(1, 1))
    )
)
)

```

[10/09 12:23:01 d2.data.build]: Removed 0 images with no usable annotations. 69 image s left.

[10/09 12:23:01 d2.data.build]: Distribution of instances among all 2 categories:

category	#instances	category	#instances
daisy	52	dandelion	29
total	81		

[10/09 12:23:01 d2.data.dataset_mapper]: [DatasetMapper] Augmentations used in traini ng: [ResizeShortestEdge(short_edge_length=(640, 672, 704, 736, 768, 800), max_size=13 33, sample_style='choice'), RandomFlip()]

[10/09 12:23:01 d2.data.build]: Using training sampler TrainingSampler

[10/09 12:23:01 d2.data.common]: Serializing 69 elements to byte tensors and concaten ating them all ...

[10/09 12:23:01 d2.data.common]: Serialized dataset takes 0.04 MiB

Skip loading parameter 'roi_heads.box_predictor.cls_score.weight' to the model due to incompatible shapes: (81, 1024) in the checkpoint but (3, 1024) in the model! You might want to double check if this is expected.

Skip loading parameter 'roi_heads.box_predictor.cls_score.bias' to the model due to incompatible shapes: (81,) in the checkpoint but (3,) in the model! You might want to double check if this is expected.

Skip loading parameter 'roi_heads.box_predictor.bbox_pred.weight' to the model due to incompatible shapes: (320, 1024) in the checkpoint but (8, 1024) in the model! You might want to double check if this is expected.

Skip loading parameter 'roi_heads.box_predictor.bbox_pred.bias' to the model due to incompatible shapes: (320,) in the checkpoint but (8,) in the model! You might want to double check if this is expected.

Skip loading parameter 'roi_heads.mask_head.predictor.weight' to the model due to incompatible shapes: (80, 256, 1, 1) in the checkpoint but (2, 256, 1, 1) in the model! You might want to double check if this is expected.

Skip loading parameter 'roi_heads.mask_head.predictor.bias' to the model due to incompatible shapes: (80,) in the checkpoint but (2,) in the model! You might want to double check if this is expected.

Some model parameters or buffers are not found in the checkpoint:

`roi_heads.box_predictor.bbox_pred.{bias, weight}`

`roi_heads.box_predictor.cls_score.{bias, weight}`

`roi_heads.mask_head.predictor.{bias, weight}`

`[10/09 12:23:01 d2.engine.train_loop]: Starting training from iteration 0`

C:\Users\shash\anaconda3\envs\detect\lib\site-packages\torch\functional.py:568: UserWarning: torch.meshgrid: in an upcoming release, it will be required to pass the indexing argument. (Triggered internally at C:\actions-runner_work\pytorch\pytorch\builder\windows\pytorch\aten\src\ATen\native\TensorShape.cpp:2228.)

`return _VF.meshgrid(tensors, **kwargs) # type: ignore[attr-defined]`

```
[10/09 12:23:19 d2.utils.events]: eta: 0:02:26 iter: 19 total_loss: 2.203 loss_cls: 1.016 loss_box_reg: 0.4728 loss_mask: 0.6943 loss_rpn_cls: 0.002908 loss_rpn_loc: 0.007563 time: 0.5924 data_time: 0.0491 lr: 1.6068e-05 max_mem: 1570M
[10/09 12:23:30 d2.utils.events]: eta: 0:02:16 iter: 39 total_loss: 1.903 loss_cls: 0.7097 loss_box_reg: 0.4722 loss_mask: 0.6723 loss_rpn_cls: 0.003308 loss_rpn_loc: 0.006597 time: 0.5775 data_time: 0.0496 lr: 3.2718e-05 max_mem: 1657M
[10/09 12:23:41 d2.utils.events]: eta: 0:02:05 iter: 59 total_loss: 1.53 loss_cls: 0.4615 loss_box_reg: 0.4609 loss_mask: 0.6329 loss_rpn_cls: 0.005835 loss_rpn_loc: 0.005716 time: 0.5588 data_time: 0.0382 lr: 4.9367e-05 max_mem: 1657M
[10/09 12:23:52 d2.utils.events]: eta: 0:01:56 iter: 79 total_loss: 1.483 loss_cls: 0.3617 loss_box_reg: 0.5195 loss_mask: 0.5873 loss_rpn_cls: 0.003574 loss_rpn_loc: 0.007116 time: 0.5649 data_time: 0.0462 lr: 6.6017e-05 max_mem: 1657M
[10/09 12:24:03 d2.utils.events]: eta: 0:01:45 iter: 99 total_loss: 1.294 loss_cls: 0.3132 loss_box_reg: 0.498 loss_mask: 0.5155 loss_rpn_cls: 0.0007185 loss_rpn_loc: 0.005508 time: 0.5585 data_time: 0.0418 lr: 8.2668e-05 max_mem: 1657M
[10/09 12:24:15 d2.utils.events]: eta: 0:01:34 iter: 119 total_loss: 1.216 loss_cls: 0.2771 loss_box_reg: 0.4954 loss_mask: 0.4422 loss_rpn_cls: 0.003772 loss_rpn_loc: 0.006196 time: 0.5628 data_time: 0.0499 lr: 9.9318e-05 max_mem: 1657M
[10/09 12:24:27 d2.utils.events]: eta: 0:01:26 iter: 139 total_loss: 1.055 loss_cls: 0.2243 loss_box_reg: 0.43 loss_mask: 0.3976 loss_rpn_cls: 0.0006666 loss_rpn_loc: 0.005985 time: 0.5700 data_time: 0.0666 lr: 0.00011597 max_mem: 1657M
[10/09 12:24:39 d2.utils.events]: eta: 0:01:17 iter: 159 total_loss: 1.13 loss_cls: 0.2383 loss_box_reg: 0.4984 loss_mask: 0.3325 loss_rpn_cls: 0.0006258 loss_rpn_loc: 0.005737 time: 0.5749 data_time: 0.0774 lr: 0.00013262 max_mem: 1657M
[10/09 12:24:54 d2.utils.events]: eta: 0:01:07 iter: 179 total_loss: 1.006 loss_cls: 0.179 loss_box_reg: 0.4991 loss_mask: 0.3058 loss_rpn_cls: 0.0009264 loss_rpn_loc: 0.006522 time: 0.5925 data_time: 0.1224 lr: 0.00014927 max_mem: 1657M
[10/09 12:25:10 d2.utils.events]: eta: 0:00:58 iter: 199 total_loss: 0.9302 loss_cls: 0.1742 loss_box_reg: 0.497 loss_mask: 0.263 loss_rpn_cls: 0.0005255 loss_rpn_loc: 0.006551 time: 0.6140 data_time: 0.1411 lr: 0.00016592 max_mem: 1657M
[10/09 12:25:26 d2.utils.events]: eta: 0:00:48 iter: 219 total_loss: 0.8432 loss_cls: 0.1191 loss_box_reg: 0.4761 loss_mask: 0.2366 loss_rpn_cls: 0.0003727 loss_rpn_loc: 0.006912 time: 0.6298 data_time: 0.1385 lr: 0.00018257 max_mem: 1657M
[10/09 12:25:42 d2.utils.events]: eta: 0:00:37 iter: 239 total_loss: 0.7571 loss_cls: 0.1308 loss_box_reg: 0.3765 loss_mask: 0.2066 loss_rpn_cls: 0.0007384 loss_rpn_loc: 0.008643 time: 0.6459 data_time: 0.1455 lr: 0.00019922 max_mem: 1657M
[10/09 12:25:54 d2.utils.events]: eta: 0:00:25 iter: 259 total_loss: 0.7593 loss_cls: 0.1127 loss_box_reg: 0.3621 loss_mask: 0.1912 loss_rpn_cls: 0.0005529 loss_rpn_loc: 0.008648 time: 0.6431 data_time: 0.0889 lr: 0.00021587 max_mem: 1657M
[10/09 12:26:03 d2.utils.events]: eta: 0:00:12 iter: 279 total_loss: 0.5024 loss_cls: 0.0802 loss_box_reg: 0.2557 loss_mask: 0.1607 loss_rpn_cls: 3.492e-05 loss_rpn_loc: 0.006331 time: 0.6270 data_time: 0.0263 lr: 0.00023252 max_mem: 1657M
[10/09 12:26:12 d2.utils.events]: eta: 0:00:00 iter: 299 total_loss: 0.5447 loss_cls: 0.08526 loss_box_reg: 0.2551 loss_mask: 0.1501 loss_rpn_cls: 8.702e-05 loss_rpn_loc: 0.006882 time: 0.6141 data_time: 0.0273 lr: 0.00024917 max_mem: 1657M
[10/09 12:26:12 d2.engine.hooks]: Overall training speed: 298 iterations in 0:03:03 (0.6141 s / it)
[10/09 12:26:12 d2.engine.hooks]: Total training time: 0:03:03 (0:00:00 on hooks)
```

In [8]:

```
#testing the data
cfg.MODEL.WEIGHTS = os.path.join(cfg.OUTPUT_DIR, "model_final.pth")
cfg.MODEL.ROI_HEADS.SCORE_THRESH_TEST = 0.5
cfg.DATASETS.TEST = ("outlierstest", )
predictor = DefaultPredictor(cfg)
```

```
[10/09 12:26:19 d2.checkpoint.c2_model_loading]: Following weights matched with mode
1:
| Names in Model | Names in Checkpoint
| Shapes | |
|:-----|:-----|:-----|
| backbone.bottom_up.res2.0.conv1.* | backbone.bottom_up.res2.0.conv1.
{norm.bias, norm.running_mean, norm.running_var, norm.weight, weight} | (64,) (64,) (6
4,) (64,) (64,64,1,1) | |
| backbone.bottom_up.res2.0.conv2.* | backbone.bottom_up.res2.0.conv2.
{norm.bias, norm.running_mean, norm.running_var, norm.weight, weight} | (64,) (64,) (6
4,) (64,) (64,64,3,3) | |
| backbone.bottom_up.res2.0.conv3.* | backbone.bottom_up.res2.0.conv3.
{norm.bias, norm.running_mean, norm.running_var, norm.weight, weight} | (256,) (256,)
(256,) (256,64,1,1) | |
| backbone.bottom_up.res2.0.shortcut.* | backbone.bottom_up.res2.0.shortcu
t.{norm.bias, norm.running_mean, norm.running_var, norm.weight, weight} | (256,) (256,)
(256,) (256,64,1,1) | |
| backbone.bottom_up.res2.1.conv1.* | backbone.bottom_up.res2.1.conv1.
{norm.bias, norm.running_mean, norm.running_var, norm.weight, weight} | (64,) (64,64,1,1)
(64,) (64,256,1,1) | |
| backbone.bottom_up.res2.1.conv2.* | backbone.bottom_up.res2.1.conv2.
{norm.bias, norm.running_mean, norm.running_var, norm.weight, weight} | (64,) (64,64,3,3)
(64,) (64,64,1,1) | |
| backbone.bottom_up.res2.1.conv3.* | backbone.bottom_up.res2.1.conv3.
{norm.bias, norm.running_mean, norm.running_var, norm.weight, weight} | (256,) (256,64,1,1)
(256,) (256,64,1,1) | |
| backbone.bottom_up.res2.2.conv1.* | backbone.bottom_up.res2.2.conv1.
{norm.bias, norm.running_mean, norm.running_var, norm.weight, weight} | (64,) (64,64,1,1)
(64,) (64,256,1,1) | |
| backbone.bottom_up.res2.2.conv2.* | backbone.bottom_up.res2.2.conv2.
{norm.bias, norm.running_mean, norm.running_var, norm.weight, weight} | (64,) (64,64,3,3)
(64,) (64,64,1,1) | |
| backbone.bottom_up.res2.2.conv3.* | backbone.bottom_up.res2.2.conv3.
{norm.bias, norm.running_mean, norm.running_var, norm.weight, weight} | (256,) (256,64,1,1)
(256,) (256,64,1,1) | |
| backbone.bottom_up.res3.0.conv1.* | backbone.bottom_up.res3.0.conv1.
{norm.bias, norm.running_mean, norm.running_var, norm.weight, weight} | (128,) (128,64,1,1)
(128,) (128,256,1,1) | |
| backbone.bottom_up.res3.0.conv2.* | backbone.bottom_up.res3.0.conv2.
{norm.bias, norm.running_mean, norm.running_var, norm.weight, weight} | (128,) (128,64,3,3)
(128,) (128,128,1,1) | |
| backbone.bottom_up.res3.0.conv3.* | backbone.bottom_up.res3.0.conv3.
{norm.bias, norm.running_mean, norm.running_var, norm.weight, weight} | (512,) (512,64,1,1)
(512,) (512,128,1,1) | |
| backbone.bottom_up.res3.0.shortcut.* | backbone.bottom_up.res3.0.shortcu
t.{norm.bias, norm.running_mean, norm.running_var, norm.weight, weight} | (512,) (512,64,1,1)
(512,) (512,256,1,1) | |
| backbone.bottom_up.res3.1.conv1.* | backbone.bottom_up.res3.1.conv1.
{norm.bias, norm.running_mean, norm.running_var, norm.weight, weight} | (128,) (128,64,3,3)
(128,) (128,128,1,1) | |
| backbone.bottom_up.res3.1.conv2.* | backbone.bottom_up.res3.1.conv2.
{norm.bias, norm.running_mean, norm.running_var, norm.weight, weight} | (128,) (128,128,3,3)
(128,) (128,128,1,1) | |
| backbone.bottom_up.res3.1.conv3.* | backbone.bottom_up.res3.1.conv3.
{norm.bias, norm.running_mean, norm.running_var, norm.weight, weight} | (512,) (512,64,1,1)
(512,) (512,128,1,1) | |
| backbone.bottom_up.res3.2.conv1.* | backbone.bottom_up.res3.2.conv1.
{norm.bias, norm.running_mean, norm.running_var, norm.weight, weight} | (128,) (128,64,3,3)
(128,) (128,128,1,1) | |
```



```

(256,) (256,) (256,256,3,3) | | backbone.bottom_up.res4.4.conv3.* | backbone.bottom_up.res4.4.conv3.
| backbone.bottom_up.res4.4.conv3.* {norm.bias,norm.running_mean,norm.running_var,norm.weight,weight} | (1024,) (102
4,) (1024,) (1024,256,1,1) | | backbone.bottom_up.res4.5.conv1.* | backbone.bottom_up.res4.5.conv1.
{norm.bias,norm.running_mean,norm.running_var,norm.weight,weight} | (256,) (256,)
(256,) (256,) (256,1024,1,1) | | backbone.bottom_up.res4.5.conv2.* | backbone.bottom_up.res4.5.conv2.
{norm.bias,norm.running_mean,norm.running_var,norm.weight,weight} | (256,) (256,)
(256,) (256,) (256,256,3,3) | | backbone.bottom_up.res4.5.conv3.* | backbone.bottom_up.res4.5.conv3.
{norm.bias,norm.running_mean,norm.running_var,norm.weight,weight} | (1024,) (102
4,) (1024,) (1024,256,1,1) | | backbone.bottom_up.res5.0.conv1.* | backbone.bottom_up.res5.0.conv1.
{norm.bias,norm.running_mean,norm.running_var,norm.weight,weight} | (512,) (512,)
(512,) (512,) (512,1024,1,1) | | backbone.bottom_up.res5.0.conv2.* | backbone.bottom_up.res5.0.conv2.
{norm.bias,norm.running_mean,norm.running_var,norm.weight,weight} | (512,) (512,)
(512,) (512,) (512,512,3,3) | | backbone.bottom_up.res5.0.conv3.* | backbone.bottom_up.res5.0.conv3.
{norm.bias,norm.running_mean,norm.running_var,norm.weight,weight} | (2048,) (204
8,) (2048,) (2048,512,1,1) | | backbone.bottom_up.res5.0.shortcut.* | backbone.bottom_up.res5.0.shortcu
t.{norm.bias,norm.running_mean,norm.running_var,norm.weight,weight} | (2048,) (2048,)
(2048,) (2048,1024,1,1) | | backbone.bottom_up.res5.1.conv1.* | backbone.bottom_up.res5.1.conv1.
{norm.bias,norm.running_mean,norm.running_var,norm.weight,weight} | (512,) (512,)
(512,) (512,) (512,2048,1,1) | | backbone.bottom_up.res5.1.conv2.* | backbone.bottom_up.res5.1.conv2.
{norm.bias,norm.running_mean,norm.running_var,norm.weight,weight} | (512,) (512,)
(512,) (512,) (512,512,3,3) | | backbone.bottom_up.res5.1.conv3.* | backbone.bottom_up.res5.1.conv3.
{norm.bias,norm.running_mean,norm.running_var,norm.weight,weight} | (2048,) (204
8,) (2048,) (2048,512,1,1) | | backbone.bottom_up.res5.2.conv1.* | backbone.bottom_up.res5.2.conv1.
{norm.bias,norm.running_mean,norm.running_var,norm.weight,weight} | (512,) (512,)
(512,) (512,) (512,2048,1,1) | | backbone.bottom_up.res5.2.conv2.* | backbone.bottom_up.res5.2.conv2.
{norm.bias,norm.running_mean,norm.running_var,norm.weight,weight} | (512,) (512,)
(512,) (512,) (512,512,3,3) | | backbone.bottom_up.res5.2.conv3.* | backbone.bottom_up.res5.2.conv3.
{norm.bias,norm.running_mean,norm.running_var,norm.weight,weight} | (2048,) (204
8,) (2048,) (2048,512,1,1) | | backbone.bottom_up.stem.conv1.* | backbone.bottom_up.stem.conv1.{no
rm.bias,norm.running_mean,norm.running_var,norm.weight,weight} | (64,) (64,)
(64,) (64,3,7,7) | | backbone.fpn_lateral2.* | backbone.fpn_lateral2.{bias,weigh
t} | (256,) (256,25
6,1,1) | | backbone.fpn_lateral3.* | backbone.fpn_lateral3.{bias,weigh
t} | (256,) (256,51
2,1,1) | | backbone.fpn_lateral4.* | backbone.fpn_lateral4.{bias,weigh
t} | (256,) (256,102
4,1,1) | | backbone.fpn_lateral5.* | backbone.fpn_lateral5.{bias,weigh
t} | (256,) (256,204
8,1,1) | | backbone.fpn_output2.* | backbone.fpn_output2.{bias,weigh
t} | (256,) (256,25

```

```

6,3,3) | backbone.fpn_output3.* | backbone.fpn_output3.{bias,weigh
| backbone.fpn_output3.* | (256,) (256,25
t} | (256,) (256,25
6,3,3) | backbone.fpn_output4.* | backbone.fpn_output4.{bias,weigh
| backbone.fpn_output4.* | (256,) (256,25
t} | (256,) (256,25
6,3,3) | backbone.fpn_output5.* | backbone.fpn_output5.{bias,weigh
| backbone.fpn_output5.* | (256,) (256,25
t} | (256,) (256,25
6,3,3) | proposal_generator.rpn_head.anchor_deltas.* | proposal_generator.rpn_head.ancho
| proposal_generator.rpn_head.anchor_deltas.* | (12,) (12,256,
r_deltas.{bias,weight} | (12,) (12,256,
1,1) | proposal_generator.rpn_head.conv.* | proposal_generator.rpn_head.conv.
| proposal_generator.rpn_head.conv.* | (256,) (256,25
{bias,weight} | (256,) (256,25
6,3,3) | proposal_generator.rpn_head.objectness_logits.* | proposal_generator.rpn_head.objec
| proposal_generator.rpn_head.objectness_logits.* | (3,) (3,256,1,
tness_logits.{bias,weight} | (3,) (3,256,1,
1) | roi_heads.box_head.fc1.* | roi_heads.box_head.fc1.{bias,weig
| roi_heads.box_head.fc1.* | (1024,) (1024,1
ht} | (1024,) (1024,1
2544) | roi_heads.box_head.fc2.* | roi_heads.box_head.fc2.{bias,weig
| roi_heads.box_head.fc2.* | (1024,) (1024,1
ht} | (1024,) (1024,1
024) | roi_heads.box_predictor.bbox_pred.* | roi_heads.box_predictor.bbox_pre
| roi_heads.box_predictor.bbox_pred.* | (8,) (8,1024
d.{bias,weight} | (8,) (8,1024
| roi_heads.box_predictor.cls_score.* | roi_heads.box_predictor.cls_scor
| roi_heads.box_predictor.cls_score.* | (3,) (3,1024
e.{bias,weight} | (3,) (3,1024
| roi_heads.mask_head.deconv.* | roi_heads.mask_head.deconv.{bias,
| roi_heads.mask_head.deconv.* | (256,) (256,25
weight} | (256,) (256,25
6,2,2) | roi_heads.mask_head.mask_fcn1.* | roi_heads.mask_head.mask_fcn1.{bi
| roi_heads.mask_head.mask_fcn1.* | (256,) (256,25
as,weight} | (256,) (256,25
6,3,3) | roi_heads.mask_head.mask_fcn2.* | roi_heads.mask_head.mask_fcn2.{bi
| roi_heads.mask_head.mask_fcn2.* | (256,) (256,25
as,weight} | (256,) (256,25
6,3,3) | roi_heads.mask_head.mask_fcn3.* | roi_heads.mask_head.mask_fcn3.{bi
| roi_heads.mask_head.mask_fcn3.* | (256,) (256,25
as,weight} | (256,) (256,25
6,3,3) | roi_heads.mask_head.mask_fcn4.* | roi_heads.mask_head.mask_fcn4.{bi
| roi_heads.mask_head.mask_fcn4.* | (256,) (256,25
as,weight} | (256,) (256,25
6,3,3) | roi_heads.mask_head.predictor.* | roi_heads.mask_head.predictor.{bi
| roi_heads.mask_head.predictor.* | (2,) (2,256,1,
as,weight} | (2,) (2,256,1,
1) | (2,) (2,256,1,

```

```
In [10]: ##for test data
test_dataset_dicts = get_data_dicts('test', classes)
```

```
In [11]: ##showing examples of the code on testset##
from matplotlib import pyplot as plt
import cv2

my_dataset_train_metadata = microcontroller_metadata
dataset_dicts = test_dataset_dicts
```

```
import random

from detectron2.utils.visualizer import Visualizer

for d in random.sample(dataset_dicts, 3):
    img = cv2.imread(d["file_name"])
    visualizer = Visualizer(img[:, :, ::-1], metadata=my_dataset_train_metadata, scale=1.2)
    vis = visualizer.draw_dataset_dict(d)
    vis=vis.get_image()[:, :, ::-1]
    #print(vis)
    plt.imshow(vis)
    plt.show()
```



